

# GARMA DKR Revamp

Caden Hewlett

2025-07-09

## Conditional Density of the Response: `dmeangamma`

Recall that the mean parameterization of the conditional distribution of the gamma-distributed response variable  $y_t$  is given in terms of the shape  $\alpha$  and mean  $\mu_t$  as:

$$f_{Y_t}(y_t | \underline{\mathbf{H}}_t) = \frac{1}{\Gamma(\alpha)} \left( \frac{\alpha}{\mu_t} \right)^\alpha y_t^{\alpha-1} \exp \left( - \frac{\alpha y_t}{\mu_t} \right)$$

To implement this in R, we first compute the log-density;

$$\log(f_{Y_t}(y_t | \underline{\mathbf{H}}_t)) = \alpha \left[ \log(\alpha) + \log(y_t) - \log(\mu_t) \right] - \log(y_t) - \log(\Gamma(\alpha)) - \frac{\alpha y_t}{\mu_t}$$

We can then write up an R function for this. The base condition for  $\mu_t$  is similar to the base parameterization in `dgamma`, wherein `scale = 1` by default.

Recall that we arrived at this parameterization by letting the scale  $\beta = \mu_t/\alpha$ . Thus, to match the `dgamma` defaults where  $\beta = 1$  we allow  $1 = \mu_t/\alpha$  and hence  $\mu_t = \alpha$  is the default functionality.

This should allow `dmeangamma(yt, shape) == dgamma(x, shape)` when only the shape is specified.

```
dmeangamma <- function(yt, shape, mu = shape, log = FALSE){  
  # input validation  
  if (any(shape < 0)) {  
    stop("Shape Parameter must be Positive!")  
  }  
  if (any((mu / shape) <= 0)) {  
    stop("Mean Parameter must be Positive!")  
  }  
  # compute log density  
  log_density <- ( shape*(log(shape) + log(yt) - log(mu)) - log(yt)  
    - lgamma(shape) - (shape*yt)/mu )  
  
  # if yt not in supp(Yt)  
  log_density[yt < 0] <- -1e10  
  
  # return  
  if(log == FALSE){  
    return(exp(log_density))  
  } else{  
    return(log_density)  
  }  
}  
# matches dgamma
```

```
alpha <- runif(1)
stopifnot(0 ==
  round(dgamma(1.8, shape = alpha) - dmeangamma(1.8, shape = alpha), 10))
```

## GARMA(p,q) Function: Identity Link

Recall the form of the GARMA( $p, q$ ) model defined previously;

$$\eta_t = g(\mu_t) = \mathbf{x}_t^\top \boldsymbol{\beta} + \tau_t$$

Where

$$\tau_t = \sum_{j=1}^p \phi_j (g(y_{t-j}) - \mathbf{x}_{t-j}^\top \boldsymbol{\beta}) + \sum_{j=1}^q \theta_j (g(y_{t-j}) - \eta_{t-j}) \quad (1)$$

Now, we allow  $g(u) = u$ . We now model the mean directly since  $\eta_t = \mu_t$ , and the ARMA component  $\tau_t$  is given by

$$\tau_t = \sum_{j=1}^p \phi_j (y_{t-j} - \tilde{y}_{t-j}) + \sum_{j=1}^q \theta_j (y_{t-j} - \mu_{t-j}), \text{ and } \mu_t = \tilde{y}_t + \tau_t$$

### Parameter Object

Recall our construction of  $\Xi = \{\mathbf{P}, \boldsymbol{\phi}, \boldsymbol{\theta}, \alpha\}$  which simplifies to  $\Xi = \{\boldsymbol{\beta}, \boldsymbol{\phi}, \boldsymbol{\theta}, \alpha\}$  in the general regression case.

The form of a parameter object  $\Xi$  is:

```
Xi_ex <- list(
  matrix(c(1, 2, 1.1, 0.9), nrow = 1), # k by 4 DKR parameter matrix
  c(0.1, 0.3), # AR parameters
  c(-0.4), # MA parameters
  2.1 # shape estimate
)
```

However, in order to pass to an optimizer we must present it in the form of a vector, i.e.

$$\Xi = \left\langle \underbrace{\beta_{01}, \beta_{11}, \delta_1, \sigma_1, \dots, \sigma_k}_{4k \text{ Kernel Regression Terms}}, \underbrace{\phi_1, \dots, \phi_p}_p, \underbrace{\theta_1, \dots, \theta_q}_q, \underbrace{\alpha}_{\text{Shape}} \right\rangle$$

We can quickly build a function to transform a vector from an optimizer into a parameter list;

```
par_to_xi <- function(P, k, p, q){
  # build empty list
  Xi <- vector("list", 4)
  # get regression parameters
  Xi[[1]] <- matrix(P[1:(4*k)], nrow = k, ncol = 4, byrow = TRUE)
  # check for AR params
  Xi[[2]] <- if (p > 0) P[(4*k + 1):(4*k + p)] else 0
  # ibid for MA params
  Xi[[3]] <- if (q > 0) P[(length(P) - q):(length(P) - 1)] else 0
  # then alpha is the last entry
  Xi[[4]] <- P[length(P)]
  # return list
  return(Xi)
}
```

This logic will be used in the prediction function later on. Further, the information object  $\mathbf{H}_t$  contains all previous response values, all covariates and the previous mean estimates. This is captured in the arguments of the state-space update equation.

## State-Space Update Equation: `compute_mu`

We assume in the code to follow that  $y[1] = y_1$ .

```
compute_mu <- function(t, y, y_tilde, mu, Xi, g = I){  
  
  # ----- #  
  # --- AR Piece --- #  
  # ----- #  
  phi <- Xi[[2]]  
  if (length(phi) > 0) {  
    idx_ar <- t - seq_along(phi)  
    term_ar <- ifelse(idx_ar > 0,  
                      phi * (g(y[idx_ar]) - y_tilde[idx_ar]),  
                      0)  
  } else {  
    term_ar <- 0  
  }  
  
  # ----- #  
  # --- MA Piece --- #  
  # ----- #  
  theta <- Xi[[3]]  
  if (length(theta) > 0) {  
    idx_ma <- t - seq_along(theta)  
    term_ma <- ifelse(idx_ma > 0,  
                      theta * (g(y[idx_ma]) - g(mu[idx_ma])),  
                      0)  
  } else {  
    term_ma <- 0  
  }  
  
  tau <- sum(term_ar) + sum(term_ma)  
  tau <- ifelse(is.na(tau), yes = 0, no = tau)  
  mu <- y_tilde[t] + tau  
  
  return(mu)  
}
```

## Worked Test Example

We verify our function with a toy example.

Consider the time points  $t \in \{1, 2, 3\}$ . The observed streamflow for these time points is  $y_t = \langle 10, 11, 12 \rangle$ . Suppose our DKR function has fit  $\tilde{y}_t = \langle 9, 10, 11 \rangle$ . We are currently fitting a  $(p, q) = (1, 1)$  model with  $\hat{\phi}_1 = 0.5$  and  $\hat{\theta}_1 = 0.3$ .

Since  $\tilde{y}_1 = 9$  is the first observation,  $\tau_1 = 0$  and thus  $\mu_1 = \tilde{y}_1 + \tau_1 = 9$ . From this information, we can compute  $\mu_2$  as follows:

$$\begin{aligned}\mu_2 &= \tilde{y}_2 + \tau_2 \\ &= \tilde{y}_2 + \sum_{j=1}^1 \hat{\phi}_j (y_{2-j} - \tilde{y}_{2-j}) + \sum_{j=1}^1 \hat{\theta}_j (y_{2-j} - \mu_{2-j}) \\ &= \tilde{y}_2 + \hat{\phi}_1 (y_1 - \tilde{y}_1) + \hat{\theta}_1 (y_1 - \mu_1) \\ &= 10 + (0.5)(10 - 9) + 0.3(10 - 9) \\ &= 10.8\end{aligned}$$

We would expect the state space update function to behave in an identical manner.

```
y <- c(10, 11)
ytilde <- c( 9, 10)
Xi <- list(NULL, c(0.5), c(0.3))
# compute first mu
mu_1 <- compute_mu(
  t = 1,
  y = y,
  y_tilde = ytilde,
  mu = NA,
  Xi = Xi
)
mu_2 <- compute_mu(
  t = 2,
  y = y,
  y_tilde = ytilde,
  mu = c(mu_1),
  Xi = Xi
)
test_that("toy example reproduces manual calculation", {
  expect_equal(mu_1, 9.0, tolerance = 1e-12)
  expect_equal(mu_2, 10.8, tolerance = 1e-12)
})
```

## Test passed

As the above implementation shows, each new  $\mu_t$  relies on prior means up to lag  $\max\{p, q\}$ . Thus, we cannot linearly apply the state space update function to our computed  $\tilde{y}_t$  values. Using a `for` loop is also inefficient for the scale of our data set (since in each iteration of the optimization we are computing 10,000 or more  $\mu_{ts}$ ,  $\tilde{y}_{ts}$  and so forth.)

## C Implementation

We can translate the code to C to increase performance speed by orders of magnitude.

```
library(Rcpp)
Rcpp::cppFunction('
    Rcpp::NumericVector compute_mu_vec(const Rcpp::NumericVector& y,
                                      const Rcpp::NumericVector& y_tilde,
                                      const Rcpp::NumericVector& phi,
                                      const Rcpp::NumericVector& theta) {

        const int    n = y.size();
        const int    p = phi.size();
        const int    q = theta.size();
        Rcpp::NumericVector mu(n), tau(n);

        for (int t = 0; t < n; ++t) {

            // ----- AR part:
            double ar = 0.0;
            for (int j = 0; j < p; ++j) {
                int idx = t - j - 1;
                if (idx >= 0)
                    ar += phi[j] * (y[idx] - y_tilde[idx]);
            }

            // ----- MA part:
            double ma = 0.0;
            for (int j = 0; j < q; ++j) {
                int idx = t - j - 1;
                if (idx >= 0)
                    ma += theta[j] * (y[idx] - mu[idx]);
            }

            tau[t] = ar + ma;
            mu[t] = y_tilde[t] + tau[t];
        }
        return mu;
    }')
```

We can verify this using our previously-verified baseline;

```
test_that("Rcpp example reproduces manual calculation", {
    expect_equal(compute_mu_vec(
        y = y,
        y_tilde = ytilde,
        phi = Xi[[2]],
        theta = Xi[[3]]
    ),
    c(9.0, 10.8),
    tolerance = 1e-12)
})
```

## Test passed

## Worked Test Example

Further, we can check that it handles the  $\max\{p, q\} \geq 2$  correctly by using our example, recalling  $y_3 = 12$ ,  $\tilde{y}_3 = 11$ ,  $p = 2$  with  $\hat{\phi}_2 = -0.15$ , to find  $\mu_3$ .

$$\begin{aligned}\mu_3 &= \tilde{y}_3 + \tau_3 \\ &= \tilde{y}_3 + \sum_{j=1}^2 \hat{\phi}_j(y_{3-j} - \tilde{y}_{3-j}) + \sum_{j=1}^1 \hat{\theta}_j(y_{3-j} - \mu_{3-j}) \\ &= \tilde{y}_3 + \hat{\phi}_1(y_2 - \tilde{y}_2) + \hat{\phi}_2(y_1 - \tilde{y}_1) + \hat{\theta}_1(y_2 - \mu_2) \\ &= 11 + 0.5(11 - 10) - 0.15(10 - 9) + 0.3(11 - 10.8) \\ &= 11.41\end{aligned}$$

```
compute_mu_vec(  
  y = c(10, 11, 12),  
  y_tilde = c(9, 10, 11),  
  phi = c(0.5, -0.15),  
  theta = 0.3  
)
```

```
## [1] 9.00 10.80 11.41
```

So, the Rcpp version is working for ARMA errors beyond (1,1).

**Runtime Comparison** Now, let's check the run time of the Rcpp version relative to the pure-R baseline for  $n = 10,000$  random observations and  $p = q = 1$ . We will use `microbenchmark` to compare runtimes.

```
set.seed(1928)
n <- 1e4
y <- rnorm(n)
ytilde <- y + arima.sim(n, model = list(ar = c(0.5), ma = c(0.3)), sd = 0.1)
phi <- 0.5; theta <- 0.3

timing <- microbenchmark(
  Rcpp = compute_mu_vec(y, ytilde, phi, theta),
  R = {
    mu <- numeric(n)
    for (t in seq_len(n))
      mu[t] <- compute_mu(t, y, ytilde, mu, list(NULL, phi, theta))
  },
  times = 10
)
```

The Rcpp implementation is exponentially faster than the pure-R alternative.

```
## Unit: microseconds
## expr      min       lq      mean   median      uq      max  neval  cld
## Rcpp    142.0    157.5    165.85    164.9    166.5    220.6     10    a
##      R 587272.7 608503.1 693408.70 676735.7 797332.5 833373.2     10    b
```



## Response Prediction Function: `predict_response_garma`

In addition to storing the different parameter objects in  $\Xi$ , we implement a new list called `configs` which is intended to decrease the bulkiness of the `predict_response` function call.

Initially, we had to declare `p`, `kernel_type`, and `use_log`. In the Gamma-GARMA version, we also have `q` and potentially `g` to pass to the function.

Further, `k` was originally inferred from the length of the `params` object, since the parameter object only contained DKR kernel parameters and AR parameters. However, the inference is much less straightforward now (as we saw with the `par_to_xi` function) and hence `k` must also be provided.

The `configs` list seeks to remedy this problem. The default construction of this object is a named list of the following structure.

```
configs_default <- list(  
  k = 1, # number of windows  
  p = 0, # AR order  
  q = 0, # MA order  
  g = I, # identity link  
  kernel_type = "gamma",  
  use_log = FALSE # log-transform y and y_tilde  
)
```

This change allows the function signature to be dramatically simplified.

```
library(DKR)  
predict_response_garma <- function(xt, zt, yt, params, configs) {  
  # default config is no AR, no MA  
  configs_default <- list(  
    k = max(1, length(params)/4 - 1),  
    p = 0, q = 0, g = I,  
    kernel_type = "gamma",  
    use_log = FALSE  
  )  
  # modify by user configuration  
  configs <- modifyList(configs_default, configs)  
  # then extract relevant hyper-parameters  
  k <- configs$k  
  p <- configs$p  
  q <- configs$q  
  # ----- #  
  # --- DKR --- #  
  # ----- #  
  pars <- as.data.frame(matrix(params[1:(4*k)], nrow = k, ncol = 4, byrow = TRUE))  
  colnames(pars) <- c("beta0", "beta1", "logdelta", "logsigma")  
  # build kernels for each set of parameters  
  kernels <- sapply(1:k,  
    function(i) build_kernel(pars$logdelta[i], pars$logsigma[i],  
                             type = configs$kernel_type)  
  )  
  # convert kernel to list if there is only one kernel  
  if (k == 1)  
    kernels <- list(kernels)  
  # zt convolution
```

```

zt_conv <- sapply(1:k, function(i)
  convolve_kernel(zt, kernels[[i]]), simplify = "matrix")
# compute beta for each kernel
beta_z_conv <- sweep(zt_conv, 2, pars$beta1, `*`)
beta <- beta_z_conv + matrix(
  pars$beta0,
  nrow = nrow(beta_z_conv),
  ncol = length(pars$beta0),
  byrow = TRUE
)
# convolve xt with kernels
xt_conv <- sapply(
  1:k,
  function(i) convolve_kernel(xt, kernels[[i]]),
  simplify = "matrix"
)
# compute kernel contributions
yt_tilde <- rowSums(xt_conv * beta)
# ----- #
# --- ARMA Error --- #
# ----- #
# check for AR params
phi <- if (p > 0) params[(4*k + 1):(4*k + p)] else 0
# ibid for MA params
theta <- if (q > 0) params[(length(params) - q):(length(params) - 1)] else 0
# then alpha is the last entry (used in log-likelihood)
alpha <- params[length(params)]
# use our mu_t computation function
mu_t <- compute_mu_vec(y = yt, y_tilde = yt_tilde, phi, theta)
# return results
list(
  mu_hat = mu_t,
  yt_tilde = yt_tilde,
  tau_t = mu_t - yt_tilde,
  beta = beta
)
}

```

## Manual Computation, End-to-End

Now that we have a complete model prediction function, we must manually verify it.

For our toy example, suppose that we have a simple one Gamma Kernel DKR model with ARMA(1, 1) errors. The parameters of this model are given by:

$$\begin{aligned}\mathbf{P} &= [\beta_0, \beta_1, \delta, \sigma] = [0.5, -0.15, 1, 0.25] \\ \phi &= \langle \phi_1 \rangle = \langle 0.45 \rangle \\ \theta &= \langle \theta_1 \rangle = \langle 0.3 \rangle\end{aligned}$$

Thus, the kernel  $\kappa$  is given by

$$\kappa_i = \frac{1}{\Gamma(\alpha)} \left( \Gamma(\alpha, i/\theta) - \Gamma(\alpha, (i+1)/\theta) \right) = \frac{1}{\Gamma(1)} \left( \Gamma(1, 4i) - \Gamma(1, 4(i+1)) \right)$$

Further, the lag range for shape  $\alpha$  and rate  $\lambda = 1/\theta$  is given by

$$\ell \in \left[ \left\lfloor \frac{\gamma^{-1}(\alpha, 10^{-3})}{\lambda} \right\rfloor, \left\lceil \frac{\gamma^{-1}(\alpha, 1 - 10^{-3})}{\lambda} \right\rceil \right]$$

Thus, the general equation for our discretized gamma kernel would be:

$$\kappa_i = \left\{ \frac{1}{\Gamma(\delta_i)} \left( \Gamma(\delta_i, \ell/\sigma_i) - \Gamma(\delta_i, (\ell+1)/\sigma_i) \right) \mid \ell \in \left[ \left\lfloor \gamma^{-1}(\delta_i, 10^{-3})\sigma_i \right\rfloor, \left\lceil \gamma^{-1}(\delta_i, 1 - 10^{-3})\sigma_i \right\rceil \right] \right\}$$

Where the range for  $\ell$  is computed via Best and Roberts (1975, source ) gamma quantile algorithm used in `qgamma`, where  $\gamma^{-1}(\dots)$  is the inverse lower incomplete gamma function.

The following computes this algorithm in Python.

```
import math
from scipy.special import gammaincinv
# gamma quantile
def gamma_quantile_function(shape, rate, p):
    return gammaincinv(shape, p) / rate
# example
range(
    math.floor(gamma_quantile_function(1, 1/0.25, 0.001)),
    math.ceil(gamma_quantile_function(1, 1/0.25, 1-0.001))
)
```

```
## range(0, 2)
```

As we see from the above,  $\ell \in [0, 2]$  in our parameter case. Thus the discretized kernel is given by

$$\kappa = \left\{ \Gamma(1, 4\ell) - \Gamma(1, 4(\ell+1)) \mid \ell \in [0, 2] \right\}$$

Which we can compute as follows:

```
library(expint)
sapply(0:(2-1), function(ell){
  gammainc(1, 4*ell)-gammainc(1, 4*(1 + ell))
})
```

```
## [1] 0.98168436 0.01798018
```

```
DKR::build_kernel(log(1), log(0.25), "gamma")
```

```
## [1] 0.98201379 0.01798621
```

Notice that in order to match our current DKR code, we have  $\ell \in [0, 2)$ . It's really a matter of semantics where we consider the last lag to be.

Now that we have computed the kernel, we let our data are given by

$$\begin{bmatrix} \mathbf{y}_T & \mathbf{x}_T & \mathbf{z}_T \end{bmatrix} = \begin{bmatrix} 2 & 1 & 1 \\ 3 & 2 & -1 \\ 4 & 3 & 1 \end{bmatrix}$$

Suppose we wish to estimate  $\mu_T$  in our complete construction.

We first compute  $\tilde{y}_T$ , recalling that  $k = 1$  hence the regression component for  $t = 1$  is given by:

$$\tilde{y}_1 = (\beta_0 + \beta_1(z_1\kappa_1))(x_1\kappa_1) = (0.50 - 0.15(1 \cdot 0.982))(1 \cdot 0.982) = 0.346$$

Similarly,

$$\begin{aligned} \tilde{y}_2 &= (\beta_0 + \beta_1(z_2\kappa_1 + z_1\kappa_2))(x_2\kappa_1 + x_1\kappa_2) \\ &= (0.50 - 0.15(-1 \cdot 0.982 + 1 \cdot 0.018))(2 \cdot 0.982 + 1 \cdot 0.018) \\ &= 1.277 \end{aligned}$$

And, recalling  $\text{length}(\kappa) = 2$ ,

$$\begin{aligned} \tilde{y}_3 &= (\beta_0 + \beta_1(z_3\kappa_1 + z_2\kappa_2))(x_3\kappa_1 + x_2\kappa_2) \\ &= (0.50 - 0.15(1 \cdot 0.982 - 1 \cdot 0.018))(3 \cdot 0.982 + 2 \cdot 0.018) \\ &= 1.0598 \end{aligned}$$

The manual computation correctly aligns with both the original `DKR::predict_target` and the new implementation:

```
##           [,1]      [,2]      [,3]
## Old 0.3463542 1.277614 1.059795
## New 0.3463542 1.277614 1.059795
```

Now, we compute  $\tau_t$  and  $\mu_t$  for each iterate, recalling  $\phi = 0.45$  and  $\theta = 0.3$ . By construction,  $\tau_1 = 0$ .

$$\begin{aligned} \tau_2 &= \phi(y_1 - \tilde{y}_1) + \theta(y_1 - \mu_1) = 0.45(2 - 0.346) + 0.3(2 - 0.346) = 1.2405 \\ \mu_2 &= \tilde{y}_2 + \tau_2 = 1.2776 + 1.2405 = \boxed{2.5181} \end{aligned}$$

Finally,  $\mu_3$  can be found via the following (noticing the MA component in effect)

$$\begin{aligned} \tau_3 &= 0.45(3 - 1.2776) + 0.3(3 - 2.5181) = 0.91965 \\ \mu_3 &= \tilde{y}_3 + \tau_3 = 1.0598 + 0.9197 = \boxed{1.9795} \end{aligned}$$

Thus, our manually computed  $\mu$  values are  $\mu_T = \{0.346, 2.5181, 1.9795\}$ . We can now *finally* verify our `predict_response_gamma` function.

```
library(DKR)
computed <- predict_response_gamma(
  xt = x,
  zt = z,
```

```

yt = y,
params = c(params, 0.45, 0.30, 1),
configs = list(k=1, p=1, q=1)
)
round(computed$mu_hat, 4)

```

```
## [1] 0.3464 2.5178 1.9795
```

Everything matches up as expected (with a bit of rounding error)!

## Log-Likelihood and Error Function

The log-likelihood of this parameter set can then be found with our previously-defined `dmeangamma` function and  $\hat{\alpha} = 1$  for simplicity.

```
sum(dmeangamma(yt = y, shape = 1, mu = computed$mu_hat, log = TRUE))
```

```
## [1] -9.532589
```

Thus, for our optimization, we can construct the error function as

```

error <- function(xt, zt, yt, params, configs) {
  # predict response using object parameters
  preds <- predict_response_gamma(xt, zt, yt, params, configs)
  # check for infinite or NA fits
  if (any(!is.finite(preds$mu_hat)) | any(is.na(preds$mu_hat))) {
    return(-1e10)
  } else{
    # otherwise return the log-likelihood using our density function
    sum(dmeangamma(
      yt = yt,
      shape = params[length(params)], # alpha is always the last entry of Xi
      mu = preds$mu_hat,
      log = TRUE
    ))
  }
}
error(x,z,y, params = c(params, 0.45, 0.30, 1),
      configs = list(k=1, p=1, q=1))

```

```
## [1] -9.532589
```